

KAIJIE ZHU

☎ +1-820-587-3174 ✉ kaijiezhu@ucsb.edu 📄 Google Scholar 🌐 Website

Education

University of California, Santa Barbara

Ph.D, Computer Science, GPA: 4.00/4.00, Advisors: [William Wang](#), [Wenbo Guo](#)

Sep. 2024 – June 2028 (expected)

California, US

Institute of Automation, Chinese Academy of Sciences

Master, Computer Science, GPA: 3.86/4.00

Sep. 2021 – June 2024

Beijing, China

Huazhong University of Science and Technology

Bachelor, ACM Class in Computer Science, GPA: 3.95/4.00

Sep. 2017 – June 2021

Wuhan, Hubei, China

Research Interests

- **Agent:** Evaluating (propose new benchmarks), improving (via SFT and RL), and securing (preventing indirect prompt injection) agent for tool use.

Selected Publications

Reasoning & Agent

TermiGen: High-Fidelity Environment and Robust Trajectory Synthesis for Terminal Agents. Submitted to NeurIPS 2026, HuggingFace Daily Paper #2

[Kaijie Zhu](#), Yuzhou Nie, Yijiang Li, Yiming Huang, Jialian Wu, Jiang Liu, Ximeng Sun, Zhenfei Yin, Lun Wang, Zicheng Liu, Emad Barsoum, William Yang Wang, Wenbo Guo

rePIRL: Learn PRM with Inverse RL for LLM Reasoning. ICML 2026

Xian Wu*, [Kaijie Zhu](#)*, Ying Zhang, Lun Wang, Wenbo Guo

DevOps-Gym: Benchmarking AI Agents in Software DevOps Cycle. ICLR 2026

Yuheng Tang*, [Kaijie Zhu](#)*, Bonan Ruan, Chuqi Zhang, Michael Yang, Hongwei Li, Suyue Guo, Tianneng Shi, Zekun Li, Christopher Kruegel, Giovanni Vigna, Dawn Song, William Yang Wang, Lun Wang, Yangruibo Ding, Zhenkai Liang, Wenbo Guo

MELON: Indirect Prompt Injection Defense via Masked Re-execution and Tool Comparison. ICML 2025

[Kaijie Zhu](#), Xianjun Yang, Jindong Wang, Wenbo Guo, William Wang

AgentOrca: A Dual-System Framework to Evaluate Language Agents on Operational Routine and Constraint Adherence. ACL 2025

Zekun Li, Shinda Huang, Jiangtian Wang, Nathan Zhang, Antonis Antoniadis, Wenyue Hua, [Kaijie Zhu](#), Sirui Zeng, William Yang Wang, Xifeng Yan

CompeteAI: Understanding the Competition Behaviors in Large Language Model-based Agents. ICML 2024 (Oral)

Qinlin Zhao, Jindong Wang, Yixuan Zhang, Yiqiao Jin, [Kaijie Zhu](#), Hao Chen, Xing Xie

AgentReview: Exploring Peer Review Dynamics with LLM Agents. EMNLP 2024 (Oral)

Yiqiao Jin, Qinlin Zhao, Yiyang Wang, Hao Chen, [Kaijie Zhu](#), Yijia Xiao, Jindong Wang

Dynamic Evaluation of LLMs

DyVal: Graph-informed Dynamic Evaluation of Large Language Models. ICLR 2024 (Spotlight) 📄 150+

[Kaijie Zhu](#)*, Jiaao Chen*, Jindong Wang, Neil Zhenqiang Gong, Diyi Yang, Xing Xie

Dynamic Evaluation of Large Language Models by Meta Probing Agents. ICML 2024

[Kaijie Zhu](#), Jindong Wang, Qinlin Zhao, Ruochen Xu, Xing Xie

PromptBench: A Unified Library for Evaluation of Large Language Models. JMLR MLOSS 📄 ★ 2.8k

[Kaijie Zhu](#), Qinlin Zhao, Hao Chen, Jindong Wang, Xing Xie

A survey on evaluation of large language models. ACM TIST 📄 6k

Yupeng Chang, Xu Wang, Jindong Wang, Yuan Wu, [Kaijie Zhu](#), Hao Chen, Linyi Yang, Xiaoyuan Yi, Cunxiang Wang, Yidong Wang, Wei Ye, Yue Zhang, Yi Chang, Philip S Yu, Qiang Yang, Xing Xie

Adversarial Robustness

PromptRobust: Towards Evaluating the Robustness of Large Language Models on Adversarial Prompts.
CCS 2024 LAMPS  **600**

Kaijie Zhu, Jindong Wang, Jiaheng Zhou, Zichen Wang, Hao Chen, Yidong Wang, Linyi Yang, Wei Ye, Neil Zhenqiang Gong, Yue Zhang, Xing Xie

Improving Generalization of Adversarial Training via Robust Critical Fine-Tuning. **ICCV 2023**

Kaijie Zhu, Xixu Hu, Jindong Wang, Xing Xie, Ge Yang

Experience

Apple

Research Intern Advisor: [Alexander Toshev](#)

- Agent RL and on policy distillation for tool use.

March 2026 – Sep. 2026
Cupertino, CA

AMD

Part-time Research Intern Advisor: [Jialian Wu](#)

- Agent training for bash tool use: TermiGen.

Oct. 2025 – March 2026
Santa Clara, CA

Microsoft Research

Research Intern Advisors: [Sheng Zhang](#), [Hoifung Poon](#)

- Proposed a new data synthetic method and ranking loss to mitigate the label noise for process reward model training.

June 2025 – Sep. 2025
Redmond, WA

Microsoft Research Asia

Research Intern Advisors: [Jindong Wang](#), [Xing Xie](#)

- [ICCV'23] Developed a robust fine-tuning strategy to enhance the generalization ability of adversarially trained models.
- [JMLR'24] Introduced PromptBench: a benchmark to evaluate the robustness of LLMs on adversarial prompts.
- [ICLR'24 Spotlight, ICML'24] Proposed a graph-informed dynamic evaluation for LLMs in reasoning tasks to mitigate test data contamination.

March 2023 – April 2024
Beijing, China

Awards

- KAUST Rising Stars in AI Symposium 2025
- Microsoft Research Star of Tomorrow, Microsoft, 2024
- Excellent Graduate Student (Top 5%), Huazhong University of Science and Technology, 2021
- Outstanding Student (Top 5%), Huazhong University of Science and Technology, 2019
- Certified Software Professional Test (Top 1%), China Computer Federation (CCF), 2019

Projects

promptbench |   **2.8k**

- Developed a flexible evaluation pipeline for large language models as the main contributor.
- Incorporated prompt engineering, dynamic evaluation for accelerating research in LLMs.

Mar. 2023 – Current

SearchAnything |   **300+**

- Created a semantic local search tool for retrieving texts and images, powered by state-of-the-art AI models.

June 2023

Tutorials & Invited Talks

- CVPR 2025 Tutorial: Multi-modal Models Evaluation: Methods and Challenges
- AAAI 2025 Tutorial: Evaluation of LLMs: Methods and Challenges
- Microsoft Research Asia 2024 LLM Evaluation Symposium